

Diachronic Word Embeddings

200 Years of Semantic Drift, Traced Through a Word-Embedding Space

Setup

All of this notebook's code lives in the package `wolfram/DiachronicEmbeddings.wl`. Evaluate the cell below first; it loads the package, which in turn reads the small committed data products in `data/` (the 1.6 GB raw `HistWords` embeddings are not needed). Every figure and table further down is then produced by a single call to a package function.

```
Get[FileNameJoin[
  {If[$Notebooks && Quiet[NotebookDirectory[]] != $Failed, ParentDirectory[
    NotebookDirectory[]], "/Users/thiel/GitHub/DiachronicWordEmbeddings"},
  "wolfram", "DiachronicEmbeddings.wl"]]
```

deHero []




Figure. Six words, two centuries. Each panel projects one word's per-decade vector – blue in the 1800s, red in the 1990s – into a 2-D map of its own changing neighbourhood. \$5 lets you explore any of the fourteen words interactively.

Word meanings move. *Gay* once meant carefree; *broadcast* meant to scatter seed by hand; a *computer* was a person who did arithmetic. This notebook makes that motion visible and measurable. Using the **HistWords** embeddings (word2vec / SGNS trained on the Google N-grams English corpus, one model per decade from the 1800s to the 1990s), we bring two centuries of independently-trained vector spaces into a single comparable frame with **orthogonal Procrustes alignment**, then trace individual words as **trajectories** through a 2-D projection of that shared space.

The drifting paths are the heart of the post (and \$5 turns them into an interactive explorer). We then test, on the same aligned embeddings, the two *statistical laws of semantic change* of Hamilton, Leskovec & Jurafsky (2016): the law of conformity (frequent words drift more slowly) replicates cleanly, while the law of innovation (polysemous words drift faster) turns

out to be a frequency confound in this corpus – a useful reminder to control before concluding. Everything is computed in the Wolfram Language; the code and data are in the companion repository.

1. The idea: words as moving points

A **word embedding** places every word at a point in a high-dimensional space so that words appearing in similar contexts land close together. Cosine similarity between two vectors then behaves like a similarity of meaning: in a modern embedding, *king* sits near *queen* and *throne*, far from *pizza*.

Now train one embedding on text from the 1850s and another on text from the 1990s. A word whose meaning has changed will sit in a different neighbourhood on the two maps. *Gay* keeps company with *cheerful*, *merry*, *lively* in the 1850s and with *homosexual*, *lesbian*, *bisexual* in the 1990s. If we can line the two maps up, that change becomes a displacement we can draw – and across twenty decades, a path.

Takeaway. Meaning change shows up as a word *moving* relative to the rest of the vocabulary. The whole project is about measuring and drawing that motion honestly.

2. The data: HistWords decade embeddings

We use the **HistWords** embeddings of Hamilton et al.: 300-dimensional skip-gram-with-negative-sampling (SGNS) vectors, trained separately for each decade on the Google N-grams eng-a11 corpus, from the 1800s through the 1990s – twenty decade-models, each with a 100,000-word vocabulary. The vectors are L2-normalized; a word absent from a decade is stored as a zero vector and skipped. The raw download is ~1.6 GB and is not committed; a thin Python shim converts the NumPy/pickle files to plain float32 binaries, and everything downstream is Wolfram Language.

Coverage varies by word. Of the fourteen we track, *gay*, *awful*, *literally*, *nice*, *mouse* and *web* appear in all twenty decades; *queer* from the 1830s, *broadcast* from the 1850s, and *computer* at 1900 (the human-“computer” sense) then continuously from the 1940s (the machine). Gaps are drawn as dashed segments. The full word list is `deWords [ ]`.

3. Why the decades must be aligned

Each decade's embedding is trained independently, so its coordinate axes are arbitrary: two decade-maps can be rotated and reflected relative to one another even when they agree perfectly about which words are near which. (A skip-gram objective is invariant under any orthogonal transformation of the space, so nothing pins down a shared orientation.) Before a word can be tracked through time, the decades must be rotated into a common frame.

We use **orthogonal Procrustes alignment**. Let A hold the vectors of the 5,000 highest-frequency words present in every decade, for one decade, and B the same words in the 1990s reference. We seek the orthogonal R minimising

$$\operatorname{argmin}_R \|AR - B\|^2 \text{ subject to } R^T R = I$$

This has a closed form from a singular value decomposition: if $U S V^T = A^T B$, then $R = U V^T$. Because R is orthogonal it preserves every length and angle *within* a decade – it only reorients the space – so each decade's nearest neighbours are exactly as trained, while becoming comparable *across* decades. (Nearest-neighbour queries in §6 are therefore run in each decade's own native space; alignment is needed only to compare frames.)

`deAlignmentPlot[]`




Figure. Mean cosine of the 5,000 shared anchor words to their 1990s vectors, before (blue) and after (red) Procrustes alignment. Alignment lifts the curve at every decade; both necessarily reach 1.0 at the 1990s reference.

How far from comparable are the raw spaces? The rotation that best aligns the 1800s onto the 1990s is large: $\|R - I\| = 21.20$. For calibration, two unrelated 300-D spaces are related by a random rotation with $\|R - I\| \approx 24.50$, identical spaces give 0, and a perfect mirror-flip gives the maximum 34.60. So the 1800s axes begin almost as unrelated to the 1990s as random – alignment is essential, not cosmetic – and the required rotation then shrinks smoothly toward the present.

Aligning lifts the mean anchor-word cosine to the 1990s at every decade (1800s: 0.42 → 0.57; 1900s: 0.50 → 0.65). That the raw (blue) curve already rises, rather than sitting near zero, is itself informative: HistWords initialises each decade's training from the previous one, so

adjacent decades inherit a partial common orientation – but that shared frame erodes over a century, which is exactly why distant decades must be re-aligned. The numbers above come from `deAlignmentFacts[ ]`.

Takeaway. Independently trained spaces are not comparable as shipped. One orthogonal rotation per decade, fitted on stable high-frequency words, makes them comparable while leaving each decade's internal geometry untouched.

4. How to read a trajectory

For one word we collect its aligned vector in every decade it occurs, gather the nearest neighbours that anchor each era of its meaning, and project that whole local neighbourhood to two dimensions with PCA. The word's path then threads between an early-meaning cluster and a late-meaning cluster; colour runs from blue (1800s) to red (1990s). One call draws any word:

```
deTrajectoryPlot["gay"]
```




Figure. *gay*: the path lingers among *merry*, *cheerful*, *joyous*, *lively*, *gaiety* through the blue and purple decades, then swings sharply red-ward to *homosexual*, *lesbian*, *gays*, *bisexual* by the 1990s – a two-century meaning change rendered as a single sweep.

What the projection does and does not say. The 2-D layout is a *per-word* PCA of that word's own cloud, chosen to spread its motion across the page; absolute positions and the


axes themselves carry no global meaning, and 2-D distances are an approximation of the true 300-D ones. The trustworthy signal is **which neighbours sit at each end** of the path and **how far the word travels between them** – both of which come straight from the high-dimensional cosine geometry, not the projection.

Takeaway. Read a trajectory by its *endpoints and their neighbours*, not by the exact shape of the curve. The neighbours are the evidence; the curve is just their geometry.

5. Explore all fourteen words

Rather than print fourteen separate plots, the same construction is wired into one interactive `Manipulate`. Pick a word, scrub the time slider to reveal the path decade by decade, and toggle the neighbour labels. (In the live notebook this is fully interactive; the static view shows *gay* with its full path.)

deExplorer[]



A few worth visiting: *broadcast* (agricultural *sow* / *seed* / *scatter* → *radio* / *television* / *bbc*); *computer* (a *V* from human surnames through *electronic* / *circuits* to *software* / *internet*); *mouse* (keeps *rat* / *cat* but grows a second *cursor* / *click* / *keyboard* cluster); *web* (spider → *www* / *browser* / *server*); *terrific* and *awful* (the two halves of a terror-to-praise swap); and *guy* (the Guy Fawkes effigy → *fellow* / *chap* / *dude*).

6. Nearest-neighbour drift tables

The trajectories are the geometry of a simpler fact: the list of a word's nearest neighbours turns over from decade to decade. Neighbours are computed by cosine similarity in each decade's own native space. Sampled every fifty years:

```
deNeighborTable[
  {"gay", "broadcast", "awful", "computer", "guy", "terrific", "media", "queer"}]
```

word	50s	00s	50s
gay	flaunting, streamers, gayly, tasteful, voluptuous	frolicsome, daft, joyous, witty, brilliant	cheerful, bright, sweet, l
broadcast	drills, sown, sows, sowing, scatter	sown, scattered, dispersed, scatter, circulated	radio, tel broadca
awful	dreadful, awfully, terrible, fearful, solemn	dreadful, terrible, horrible, appalling, mysterious	dreadful, terrible fearful
computer	-	williams, cush, burt, dietary, barney	programmed digital electro
guy	brodie, carleton, hospital, beauchamp, hunterian	maupassant, fawkes, bevis, faux, carrington	damn, joe funny, :
terrific	tremendous, terrible, dreadful, appalling, tornado	terrible, tremendous, dreadful, frightful, horrible	terrible, roar, t fierce,
media	armenia, chaldeia, phrygia, syria, libya	otitis, parthia, medium, intima, opacities	medium, te fluids, communi
queer	funny, medley, youngster, odd, comical	funny, odd, ugly, strange, silly	funny, str ugly, s

Figure. Top-5 cosine nearest neighbours by decade. Any subset of `deWords[ ]` can be passed to `deNeighborTable`.

Takeaway. Trajectories can be cross-checked against plain word lists. When the neighbours of *broadcast* go from *sow, scatter* to *radio, television*, the drawn path is just reporting that shift.

7. Two statistical laws of semantic change

Hamilton, Leskovec & Jurafsky (2016) propose two regularities of semantic change. Both can be tested directly on the aligned embeddings, measured over the **11332** dictionary content words present in every decade (function words and archaic/foreign tokens removed via a stoplist and `WordData` membership). A word's *rate of change* is the mean cosine distance between its aligned vectors in consecutive decades; its *frequency* is proxied by its mean rank, and its *polysemy* by its number of WordNet senses. Re-deriving the laws

independently is also a good way to watch a confound masquerade as a result. (Statistics: `deLawFacts[ ]`.)

The law of conformity – holds

More frequent words change more slowly. Change-rate falls steeply with frequency (Spearman $\rho = -0.94$, slope -0.12 in log-rate). *Why:* a high-frequency word is seen in a huge range of contexts, so its vector is an average over many uses and is statistically well-anchored. Part of the strength is also methodological – rare words are estimated from little data, so their vectors are noisier and *look* like they move more – which is itself worth stating plainly.

`deConformityPlot[ ]`




Figure. Rate of semantic change vs. frequency ($-\log$ mean rank). Red markers are decile means.

The law of innovation – a cautionary tale

The published law says **more polysemous words change faster**. In this corpus the *raw* correlation comes out the wrong way: change-rate vs. polysemy has Spearman $\rho = -0.31$ – negative. The culprit is a **confound**: polysemous words tend to be frequent ($\rho = 0.34$), and frequency strongly suppresses change. In the raw comparison the frequency effect simply dominates – a textbook Simpson's paradox.

Controlling for frequency (partial Spearman correlation, holding frequency rank fixed) the spurious negative trend vanishes – but it does not turn convincingly positive either: partial $\rho = 0.01$, essentially zero. **So in the Google-N-grams eng-all corpus, with WordNet sense**

counts as the polysemy measure, the law of innovation does not replicate. That is a more honest – and more useful – result than forcing the sign: the law may need a different corpus, a better polysemy measure, or the original authors' frequency-and-polysemy-matched controls to surface.

```
deInnovationPlot[]
```




Figure. The raw polysemy trend (blue) slopes the "wrong" way only because polysemous words are frequent; the frequency-controlled residual (red) is flat. Raw Spearman -0.31, partial 0.01.

Takeaway. Conformity replicates cleanly; innovation does not survive a frequency control here. A large raw correlation is easy to get and easy to misread – always ask what else moves with your variable before naming an effect.

8. The most-drifted words – and a caveat

Ranking every stable content word by how far it moved between its 1800s and 1990s aligned vectors surfaces two phenomena at once. Some entries are genuine meaning change (*advert* verb → noun, *peer* noble → equal, *media*, *broadcast*); others are lexical obsolescence or register shift, where an archaic word (*twelvemonth*, *fourthly*, *civilities*) simply stopped occurring in the same kinds of text. A large embedding displacement means “the contexts changed” – which is necessary, but not sufficient, for “the meaning changed”. Distinguishing the two is where embeddings end and lexicography begins.

deDriftChart[25]




Figure. The 25 most-drifted stable content words, 1800s–1990s. The full ranking is in `data/drift_ranking.json`.

Takeaway. Embeddings are a powerful detector of *that* something changed about a word's usage; deciding *what* changed still needs the neighbour lists, the trajectories, and human judgement.

9. Methods and reproduction

All notebook code is in `wolfram/DiachronicEmbeddings.wl`; each output above is a single call to one of its functions, drawing from the small committed data products in `data/`. To regenerate those products from scratch (the science is pure Wolfram Language; the only Python is a NumPy/pickle format shim):

- `wolframscript -file wolfram/fetch_data.wls` - download + unzip the HistWords eng-all SGNS embeddings
- `python3 src/convert_histwords.py - .npz/.pkl → float32 binaries + vocab` (anchors, targets, pools)
- `wolframscript -file wolfram/trajectories.wls` - Procrustes alignment, projection → `trajectories.json` + `neighbors.json`
- `wolframscript -file wolfram/laws_and_drift.wls` - drift leaderboard + the two laws → `laws.json`, `laws_points.json`, `drift_ranking.json`
- `wolframscript -file wolfram/alignment.wls` - the alignment diagnostic → `alignment.json`

Repository: github.com/mthiel74/DiachronicWordEmbeddings.

10. References

- W. L. Hamilton, J. Leskovec, D. Jurafsky (2016). *Diachronic Word Embeddings Reveal Statistical Laws of Semantic Change*. ACL 2016. arXiv:1605.09096.
- HistWords project and data: <https://nlp.stanford.edu/projects/histwords/> (Stanford NLP).
- P. H. Schönemann (1966). *A generalized solution of the orthogonal Procrustes problem*. Psychometrika 31(1). (The SVD alignment used in §3.)
- T. Mikolov et al. (2013). *Distributed Representations of Words and Phrases and their Compositionality*. NeurIPS 2013. (word2vec / SGNS.)
- V. Kulkarni et al. (2015). *Statistically Significant Detection of Linguistic Change*. WWW 2015.
- J.-B. Michel et al. (2011). *Quantitative Analysis of Culture Using Millions of Digitized Books*. Science 331. (The Google N-grams corpus.)